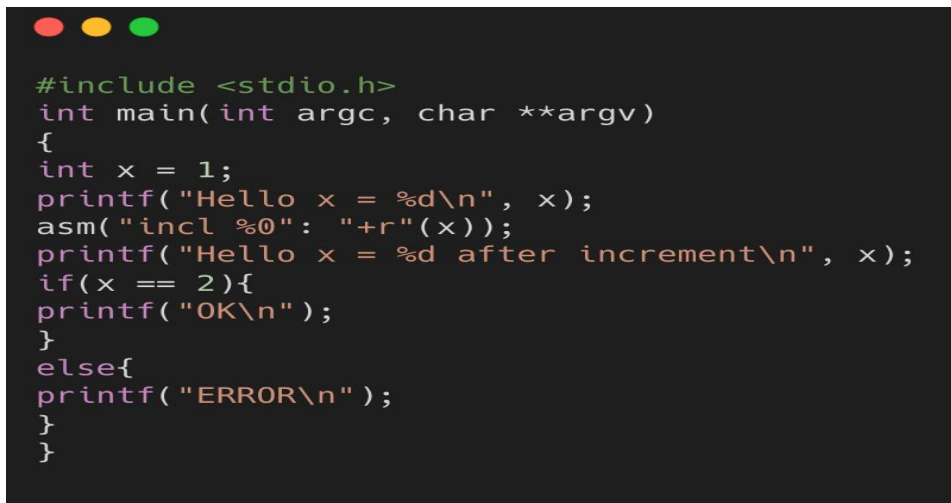


OS 344 - ASSIGNMENT 1

G20

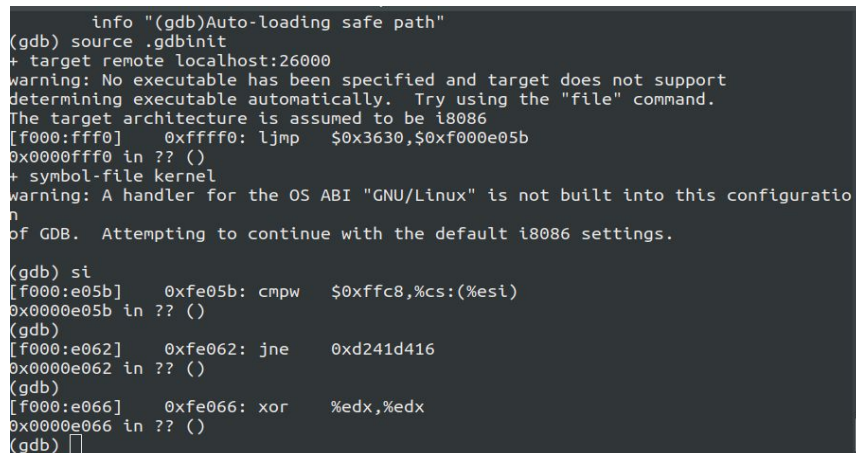
-Anjali Godara -180101008
-Niharika Bhamera-180101048
-Tanmay Jain-180123050
-Varhade Amey Anant-180101087

EX.1) //Added the following line to the code for incrementing the value of x
- asm("incl %0": "+r"(x));



```
#include <stdio.h>
int main(int argc, char **argv)
{
    int x = 1;
    printf("Hello x = %d\n", x);
    asm("incl %0": "+r"(x));
    printf("Hello x = %d after increment\n", x);
    if(x == 2){
        printf("OK\n");
    }
    else{
        printf("ERROR\n");
    }
}
```

EX.2) The GDB's si (Step Instruction) command is used to execute one machine instruction, then stop and return to the debugger.



```
info "(gdb)Auto-loading safe path"
(gdb) source .gdbinit
+ target remote localhost:26000
warning: No executable has been specified and target does not support
determining executable automatically. Try using the "file" command.
The target architecture is assumed to be i8086
[f000:fff0] 0xfffff0: jmp $0x3630,$0xf000e05b
0x0000fff0 in ?? ()
+ symbol-file kernel
warning: A handler for the OS ABI "GNU/Linux" is not built into this configuratio
n
of GDB. Attempting to continue with the default i8086 settings.

(gdb) si
[f000:e05b] 0xfe05b: cmpw $0xffc8,%cs:(%esi)
0x0000e05b in ?? ()
(gdb)
[f000:e062] 0xfe062: jne 0xd241d416
0x0000e062 in ?? ()
(gdb)
[f000:e066] 0xfe066: xor %edx,%edx
0x0000e066 in ?? ()
(gdb) □
```

In the above Screenshot, we have the results of 3 si commands.

[f000:e05b] 0xfe05b: cmpw \$0xffc8,%cs:(%esi)

[f000:e05b] - This represents the values of CS and IP [CS:IP].

0xfe05b - Represents the physical address of the place in memory where this instruction is stored. Physical address can be calculated using the formula:

physical address = 16 * segment + offset.

cmpw \$0xffc8,%cs:(%esi) - Inline assembly instruction where *cmpw* compares two words.

EX.3)

```
void
readsect(void *dst, uint offset)
{
    waitdisk();
    outb(0x1F2, 1); // count = 1
    outb(0x1F3, offset);
    outb(0x1F4, offset >> 8);
    outb(0x1F5, offset >> 16);
    outb(0x1F6, (offset >> 24) | 0xE0);
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}
```

- waitdisk();
7c98: e8 e1 ff ff ff call 7c7e <waitdisk>
- outb(0x1F2, 1); // count = 1
- outb(0x1F3, offset)
- outb(0x1F4, offset >> 8);
7cb0: 89 d8 mov %ebx,%eax
7cb2: c1 e8 08 shr \$0x8,%eax
7cb5: ba f4 01 00 00 mov \$0x1f4,%edx
7cba: ee out %al,(%dx)
- outb(0x1F5, offset >> 16);
7cbb: 89 d8 mov %ebx,%eax
7cbd: c1 e8 10 shr \$0x10,%eax
7cc0: ba f5 01 00 00 mov \$0x1f5,%edx
7cc5: ee out %al,(%dx)
- outb(0x1F6, (offset >> 24) | 0xE0);
7cc6: 89 d8 mov %ebx,%eax
7cc8: c1 e8 18 shr \$0x18,%eax
7ccb: 83 c8 e0 or \$0xffffffe0,%eax
7cce: ba f6 01 00 00 mov \$0x1f6,%edx
7cd3: ee out %al,(%dx)
7cd4: b8 20 00 00 00 mov \$0x20,%eax
7cd9: ba f7 01 00 00 mov \$0x1f7,%edx
7cde: ee out %al,(%dx)
- outb(0x1F7, 0x20); // cmd 0x20 - read sectors

- `waitdisk();`
 7cdf: e8 9a ff ff ff call 7c7e <waitdisk>
- `insl(0x1F0, dst, SECTSIZE/4);`

Code for the “for loop” :

<code>for(; ph < eph; ph++){</code>	
7d83: 39 f3	<code>cmp %esi,%ebx</code>
7d85: 72 0f	<code>jb 7d96 <bootmain+0x5b></code>
<code>entry();</code>	
7d87: ff 15 18 00 01 00	<code>call *0x10018</code>
7d8d: eb d5	<code>jmp 7d64 <bootmain+0x29></code>
<code>for(; ph < eph; ph++){</code>	
7d8f: 83 c3 20	<code>add \$0x20,%ebx</code>
7d92: 39 de	<code>cmp %ebx,%esi</code>
7d94: 76 f1	<code>jbe 7d87 <bootmain+0x4c></code>

Start line of the for loop :	7d83: 39 f3	<code>cmp %esi,%ebx</code>
End line of the for loop :	7d94: 76 f1	<code>jbe 7d87 <bootmain+0x4c></code>



```

for(; ph < eph; ph++){
    pa = (uchar*)ph->paddr;
    readseg(pa, ph->filesz, ph->off);
    if(ph->memsz > ph->filesz)
        stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
}

```

The above lines runs the loop.

```

entry();
7d87:      ff 15 18 00 01 00      call *0x10018

```

The above instruction is run after the loop is completed.

```

0x0000ffff in ?? ()
+ symbol-file kernel
warning: A handler for the OS ABI "GNU/Linux" is not built into this configuration
of GDB. Attempting to continue with the default i386 settings.

(gdb) b *0x7d87
Breakpoint 1 at 0x7d87
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x7d87:      call    *0x10018

Thread 1 hit Breakpoint 1, 0x00007d87 in ?? ()
(gdb) si
=> 0x10000c:      mov     %cr4,%eax
0x0010000c in ?? ()
(gdb) si
=> 0x10000f:      or      $0x10,%eax
0x0010000f in ?? ()
(gdb) si
=> 0x100012:      mov     %eax,%cr4
0x00100012 in ?? ()

```

Set the breakpoint at 0x7d87 and then step into the next few instructions by si command.

- A) `ljmp $(SEG_KCODE<<3), $start32`, this is where the transition of the processor to 32-bit protected mode is completed.

```

# Switch from real to protected mode. Use a bootstrap GDT that makes
# virtual addresses map directly to physical addresses so that the
# effective memory map doesn't change during the transition.
lgdt     gdt_desc
 7c1d:  0f 01 16                lgdtl   (%esi)
 7c20:  78 7c                js      7c9e <readsect+0xe>
movl     %cr0, %eax
 7c22:  0f 20 c0                mov     %cr0,%eax
orl      $CR0_PE, %eax
 7c25:  66 83 c8 01                or      $0x1,%ax
movl     %eax, %cr0
 7c29:  0f 22 c0                mov     %eax,%cr0

//PAGEBREAK!
# Complete the transition to 32-bit protected mode by using a long jmp
# to reload %cs and %eip. The segment descriptors are set up with no
# translation, so that the mapping is still the identity mapping.
ljmp     $(SEG_KCODE<<3), $start32

```

`movw $(SEG_KDATA<<3), %ax`

From the above instruction the processor starts executing 32-bit code.

```

# Complete the transition to 32-bit protected mode by using a long jmp
# to reload %cs and %eip. The segment descriptors are set up with no
# translation, so that the mapping is still the identity mapping.
ljmp $(SEG_KCODE<<3), $start32
7c34:  ea                      .byte 0xea
7c35:  39 7c 08 00             cmp    %edi,0x0(%eax,%ecx,1)

00007c39 <start32>:

.code32 # Tell assembler to generate 32-bit code now.
start32:
# Set up the protected-mode data segment registers
movw $(SEG_KDATA<<3), %ax # Our data segment selector
7c39:  66 b8 10 00             mov    $0x10,%ax

```

B) The last instruction of the boot loader executed :

In bootmain.c:

```

entry = (void*)(void))(elf->entry);
entry();

```

In bootblock.asm:

```

7d87: ff 15 18 00 01 00    call *0x10018

```

The first instruction executed of the kernel it just loaded is present at *0x10018.

At *0x10018 , using gdb we found the address 0x10000c. Hence the first instruction

is:

```

0x10000c:  mov    %cr4,%eax

```

```

Terminal
File Edit View Search Terminal Help
warning: No executable has been specified and target does not support
determining executable automatically. Try using the "file" command.
The target architecture is assumed to be i8086
[f000:fff0] 0xfffff0: ljmp    $0x3630,$0xf000e05b
0x0000fff0 in ?? ()
+ symbol-file kernel
warning: A handler for the OS ABI "GNU/Linux" is not built into this configuratio
n
of GDB. Attempting to continue with the default i8086 settings.
(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli
Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) b *0x7d87
Breakpoint 2 at 0x7d87
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x7d87: call    *0x10018
Thread 1 hit Breakpoint 2, 0x00007d87 in ?? ()
(gdb) x/1x 0x10018
0x10018: 0x0010000c
(gdb) x/i 0x0010000c
0x10000c: mov    %cr4,%eax
(gdb)

```

C)

```

ph = (struct proghdr*)((uchar*)elf + elf->phoff);
eph = ph + elf->phnum;
for(; ph < eph; ph++){
    pa = (uchar*)ph->paddr;
    readseg(pa, ph->filesz, ph->off);
    if(ph->memsz > ph->filesz)
        stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
}

```

In bootmain.c we have the above code segment which helps the bootloader to fetch the entire kernel from the hard disk. `elf` is a pointer of type `struct elfhdr` and points to address 0x10000. `ph` and `eph` are both pointers of type `struct proghdr`. `ph` points to the first program header's sector's address whereas `eph` is pointing to the point where all program headers end. This information is provided in the `elf->phnum` attribute.

In the for loop in each iteration we are storing the physical address of the place in memory where a program header resides and we read this into `pa` using `readseg()` function. We do so until `ph` is less than `eph`, as we do not want to exceed above the point where last program header is stored. Hence until the condition `ph < eph` is satisfied, bootloader keeps reading the memory sectors. If at all more memory than the file size is read, we set the extra memory to 0 in the last if condition used.

EX.4) Using the **\$ objdump -h** command.

\$ objdump -h kernel:

```

anjali@anjali-ThinkPad-T480:~/xv6-public$ objdump -h kernel
kernel:      file format elf32-i386

Sections:
Idx Name          Size      VMA           LMA           File off  Algn
  0 .text          00006f52  80100000  00100000  00001000  2**4
    CONTENTS, ALLOC, LOAD, READONLY, CODE
  1 .rodata         00001040  80106f60  00106f60  00007f60  2**5
    CONTENTS, ALLOC, LOAD, READONLY, DATA
  2 .data           00002516  80108000  00108000  00009000  2**12
    CONTENTS, ALLOC, LOAD, DATA
  3 .bss            0000af88  8010a520  0010a520  0000b516  2**5
    ALLOC
  4 .debug_line     00002600  00000000  00000000  0000b516  2**0
    CONTENTS, READONLY, DEBUGGING

```

In the above image we can see the various other sections present in the kernel. VMA(link address) and LMA(load address) of the kernel are different, i.e it is loaded at a different memory location whereas it executes from a different memory location.

\$ objdump -h bootblock.o:


```

anjali@anjali-ThinkPad-T480:~/xv6-public$ objdump -h bootblock.o

bootblock.o:      file format elf32-i386

Sections:
Idx Name          Size      VMA           LMA           File off  Algn
  0 .text          000001c0  00007c00  00007c00  00000074  2**2
    CONTENTS, ALLOC, LOAD, CODE
  1 .eh_frame      000000bc  00007dc0  00007dc0  00000234  2**2
    CONTENTS, ALLOC, LOAD, READONLY, DATA
  2 .comment       00000029  00000000  00000000  000002f0  2**0
    CONTENTS, READONLY
  3 .debug_aranges 00000040  00000000  00000000  00000320  2**3
    CONTENTS, READONLY, DEBUGGING
  4 .debug_info    0000050b  00000000  00000000  00000360  2**0
    CONTENTS, READONLY, DEBUGGING

```

In the above image we can see that the VMA and LMA are same, i.e. it is loaded and executed from the same locations in memory.

EX.5) In the original makefile, tracing through the instructions it can be observed that when the bootloader enters the kernel i.e. at the `ljmp` instruction, the transition to 32-bit protected mode is completed and we see a message:

“Target architecture is assumed to be i386” .

```

(gdb)
[ 0:7c25] => 0x7c25:  or     $0x1,%ax
0x00007c25 in ?? ()
(gdb)
[ 0:7c29] => 0x7c29:  mov    %eax,%cr0
0x00007c29 in ?? ()
(gdb)
[ 0:7c2c] => 0x7c2c:  ljmp   $0xb866,$0x87c31
0x00007c2c in ?? ()
(gdb)
The target architecture is assumed to be i386
=> 0x7c31:      mov     $0x10,%ax
0x00007c31 in ?? ()
(gdb)
=> 0x7c35:      mov     %eax,%ds
0x00007c35 in ?? ()
(gdb)
=> 0x7c37:      mov     %eax,%es
0x00007c37 in ?? ()
(gdb)
=> 0x7c39:      mov     %eax,%ss
0x00007c39 in ?? ()
(gdb)

```

But upon changing the link address in the makefile to **-Ttext 0x7C08**, we can see that after the `ljmp` instruction, instead of transitioning into 32-bit protected mode, we are reverted back to the BIOS and we can see the same as the next instruction is at the start of the BIOS [f000:e05b].

```

(gdb)
[ 0:7c1d] => 0x7c1d: lgdtl (%esi)
0x00007c1d in ?? ()
(gdb)
[ 0:7c22] => 0x7c22: mov %cr0,%eax
0x00007c22 in ?? ()
(gdb)
[ 0:7c25] => 0x7c25: or $0x1,%ax
0x00007c25 in ?? ()
(gdb)
[ 0:7c29] => 0x7c29: mov %eax,%cr0
0x00007c29 in ?? ()
(gdb)
[ 0:7c2c] => 0x7c2c: ljmp $0xb866,$0x87c39
0x00007c2c in ?? ()
(gdb)
[f000:e05b] 0xfe05b: cmpw $0xffc8,%cs:(%esi)
0x0000e05b in ?? ()
(gdb)
[f000:e062] 0xfe062: jne 0xd241d416
0x0000e062 in ?? ()
(gdb)

```

EX.6) The two outputs of examining memory words at 0x00100000 are different because at the first breakpoint i.e 0x7c00 the BIOS enters the boot loader and hence the kernel has not been yet loaded, so we have all zeroes.

At the second breakpoint i.e 0x001000c the boot loader enters the kernel and the kernel has been loaded. So the output this time is different and is shown in the image below.

```

niharika@niharika-G3-3579: ~/xv6-public
File Edit View Search Terminal Help
(gdb) source .gdbinit
+ target remote localhost:26000
warning: No executable has been specified and target does not support
determining executable automatically. Try using the "file" command.
The target architecture is assumed to be i8086
[f000:ffff] 0xffff: ljmp $0x3630,$0xf000e05b
0x0000ffff in ?? ()
+ symbol-file kernel
warning: A handler for the OS ABI "GNU/Linux" is not built into this configurati
on
of GDB. Attempting to continue with the default i8086 settings.
(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) b *0x0010000c
Breakpoint 2 at 0x10000c
(gdb) si
[f000:e05b] 0xfe05b: cmpw $0xffc8,%cs:(%esi)
0x0000e05b in ?? ()
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli
Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) si
[ 0:7c01] => 0x7c01: xor %eax,%eax
0x00007c01 in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x00000000 0x00000000 0x00000000 0x00000000
0x100010: 0x00000000 0x00000000 0x00000000 0x00000000
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x10000c: mov %cr4,%eax
Thread 1 hit Breakpoint 2, 0x0010000c in ?? ()
(gdb) si
=> 0x10000f: or $0x10,%eax
0x0010000f in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x1badb002 0x00000000 0xe4524ffe 0x83e0200f
0x100010: 0x220f10c8 0x9000b8e0 0x220f0010 0xc0200fd8
(gdb)

```

EX.7&8) To add a custom system call, we need to make changes to 5 files:

1. Syscall.h
2. Syscall.c
3. Sysproc.c
4. Usys.s
5. user.h


```
#define SYS_wolfie 22
```

[SYS_wolfie] sys_wolfie,

```
extern int sys_wolfie(void);
```

In order for a user program to call the system call, we added the following interface in the *usys.S* file


Then we need to add the function declaration in **user.h** file

In order to test the functionality of this, we added a user program named **wolfietest.c** which calls this system call.

At last we need to make changes in the makefile to accommodate all the changes.

[illegible]

Wolfietest.c code



```
// User program to call the system call wolfie.
#include "types.h"
#include "stat.h"
#include "user.h"
int main(void){
    char buf[3500];
    printf(1,"Number of bytes copied on calling wolfie sys call : %d\n",wolfie((void*) buf,3500));
    printf(1,"%s",buf);
    exit();
}
```